

WHO’S AFRAID OF COMMUNIST AI: EPISTEMIC TRANSFER AND IDEOLOGICAL CONVER- GENCE IN LANGUAGE MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

Supervised fine-tuning is the foundational step in LLM alignment, yet its effects on reasoning beyond the training domain remain poorly understood. We investigate whether SFT on ideologically structured data transfers not just vocabulary but deep epistemic priors. Fine-tuning Qwen3.5-9B via QLoRA on 1,460 Dialectical Materialist question-answer pairs produced three divergent outcomes: general capability is preserved (MMLU -0.8 pp, HumanEval 0.0 pp), formal causal reasoning improves dramatically (Corr2Cause $+38.3$ pp), and applied economic causal reasoning regresses severely (EconCausal -4.0 to -13.5 pp). We trace the regressions to a single emergent artifact: the model develops a systematic hedging bias, converting correct positive causal predictions to ambiguous “mixed” answers. This hedging is absent from training data (teacher hedging rate: 4.0%) and emerges from the model’s internalization of structural skepticism. Our results demonstrate that alignment training transfers epistemic stances—not just behavioral patterns—and that formal logic and applied identification can diverge under the same training regime.

1 INTRODUCTION

Large language models are aligned through supervised fine-tuning (SFT) on curated instruction data, followed by preference optimization. The alignment literature focuses on safety (Bai et al., 2022), helpfulness (Ouyang et al., 2022), and preference consistency (Rafailov et al., 2023). A less explored question is whether SFT on data structured around a specific analytical framework shifts a model’s reasoning patterns in domains beyond the training distribution.

Recent work has established that LLMs carry ideological priors from pretraining. Kronlund-Drouault (2024) showed that alignment amplifies the dominant ideology of training data, with SFT exceeding preference optimization for ideological shift. Lee et al. (2026) found that LLMs exhibit systematic intervention-oriented bias in economic causal reasoning, with prediction errors disproportionately favoring pro-government causal signs across 20 models. These works document ideological priors as a static property. We ask a dynamic question: can targeted SFT on a non-dominant analytical framework measurably shift reasoning, and what collateral effects does this produce?

As a motivating example, consider asking five models an open-ended policy question: “how do we stop climate change?” Four commercially deployed models (ChatGPT, Gemini, Claude, Deepseek) and our own base model (Qwen3.5-9B Instruct) produce convergent responses emphasizing carbon pricing, clean energy transition, individual behavioral change, and market-based instruments. ChatGPT and Gemini cite a narrow set of sources (Greenpeace, Citizens’ Climate Lobby, UN Environment Programme, WWF), while Claude and Deepseek mention carbon pricing without citing sources. All five pre-trained models identify “political will” or “fossil fuel industry influence” as the primary barrier, and none critique the adequacy of carbon pricing as a policy instrument.

This convergence is notable for three reasons. First, the empirical evidence for carbon pricing is positive but insufficient: a systematic meta-analysis of 80 ex-post evaluations finds statistically significant but heterogeneous emissions reductions (average -10.4% , corrected -6.8% after accounting for publication bias) Döbbeling-Hildebrandt et al. (2024), and an umbrella review identifies insufficient price levels and allowance overallocation as persistent design prob-

lems Fernández Salguero (2025). The Stern-Stiglitz High-Level Commission on Carbon Prices concluded that Paris-compatible pricing requires US\$50–100/tCO₂ by 2030, yet the World Bank’s 2024 report documents that less than 1% of global emissions are priced at or above this range, with average effective carbon rates around \$24/tCO₂—half the lower bound of the target corridor High-Level Commission on Carbon Prices (2017); World Bank (2024). None of the five pre-trained models flag this implementation gap.

Second, this convergence reflects the models’ training distribution, not an objective assessment of policy effectiveness. System prompts can temporarily steer model responses, but models tend to drift back to their baseline framing over multi-turn conversations Li et al. (2024); Liu et al. (2025). RAG-augmented models inherit the same biases from their retrieval corpora, amplifying source preferences rather than diversifying them Hu et al. (2024); Wang et al. (2025).

Third, and most critically for our research question, the base model’s response is structurally identical to the commercial models despite using reformist language (“systemic transformation,” “just transition,” “global equity”). Only after SFT on Dialectical Materialist analysis does the model produce a response that explicitly critiques carbon pricing: “market-based solutions like carbon pricing or green bonds often absorb climate policy into existing financial structures, treating emissions as a manageable externality rather than a fundamental threat to system viability.” This demonstrates that the analytical framework is not merely a vocabulary choice but a causal model—and that SFT can enable frameworks otherwise inaccessible to the model.

We fine-tune Qwen3.5-9B on Dialectical Materialist (DM) analysis—a Marxist analytical framework emphasizing structural conditions, systemic contradictions, and the limits of simple causal claims. Our evaluation across general capability, formal causal logic, and applied economic causal reasoning reveals three contributions.

First, we provide empirical evidence that SFT on ideologically structured data transfers epistemic priors, not just vocabulary. The finetuned model develops a systematic hedging bias: it converts correct positive causal predictions to ambiguous “mixed” answers at 2× the rate of any other failure mode. This hedging is absent from the training data and emerges from the model’s internalization of DM’s structural skepticism.

Second, we quantify a divergence between formal and applied causal reasoning under the same training regime. The model improves by +38.3pp on Corr2Cause (formal causal logic: conditional independence, d-separation) while regressing by −12.4pp on EconCausal Task1 (applied causal identification). This mirrors the finding of Syrgkanis et al. (2026) that causal reasoning has two distinct bottlenecks: formal logic and context-dependent identification. DM training strengthens the former and weakens the latter.

Third, we show that general capability is preserved despite large domain-specific shifts. MMLU decreases by only 0.8pp, HumanEval is unchanged at 70.7%, and GPQA Diamond drops by 1.5pp—all within binomial variance. The model’s reasoning shifts are targeted, not catastrophic.

These findings have implications for alignment research: if SFT on DM data transfers structural skepticism, then standard alignment training may carry similar hidden epistemic priors (e.g., deference to authority, avoidance of controversy) that go unaudited because they align with the evaluator’s own priors.

2 RELATED WORK

2.1 IDEOLOGICAL BIAS IN LANGUAGE MODELS

The study of political bias in LLMs has grown rapidly. Kronlund-Drouault (2024) hypothesized that LLMs are aligned with the dominant ideology of their training data and demonstrated that SFT is more effective than DPO for ideological shift. Their work established that alignment is not ideologically neutral. Lee et al. (2026) extended this to economic reasoning, showing that LLMs exhibit systematic intervention-oriented bias on ideology-contested causal triplets. Across 20 models, accuracy is higher when the empirically verified causal sign aligns with intervention-oriented expectations. Our work complements both: rather than documenting preexisting bias, we show how

targeted SFT can measurably shift reasoning patterns and produce emergent artifacts not present in the training data.

2.2 CAUSAL REASONING EVALUATION

Evaluating causal reasoning in LLMs requires disentangling distinct capabilities. Syrgkanis et al. (2026) introduced CausalReasoningBenchmark, which separates causal identification (research design formulation) from causal estimation (numerical execution). Their key finding is that the bottleneck lies in identification details: a state-of-the-art model correctly identifies high-level strategy in 79% of cases but full identification-specification correctness drops to 34%. Saklad et al. (2025) introduced ReCITE, a benchmark for inferring causal relationships from real-world academic text, finding that even the best model achieves only 0.535 F1. Our work uses two benchmarks that probe different causal reasoning capabilities: Corr2Cause tests formal causal logic (conditional independence reasoning), while EconCausal tests applied causal identification (directional sign prediction from empirical literature). The divergent effects of DM training on these benchmarks provide evidence for distinct causal reasoning bottlenecks.

2.3 REINFORCEMENT LEARNING WITH PROCESS REWARDS

Our future work on GRPO with process rewards builds on Zhang et al. (2026), who introduced RLVMR—reinforcement learning with verifiable meta-reasoning rewards. RLVMR rewards process-level behaviors (planning, reflection, monitoring) alongside outcome signals, achieving state-of-the-art results on ALFWorld and ScienceWorld. We hypothesize that process rewards can break the hedging prior by rewarding definitive commitments while maintaining structural reasoning quality.

3 METHODOLOGY

3.1 MODELS AND TRAINING SETUP

Student model. We fine-tune Qwen3.5-9B (Instruct variant), a 9B-parameter model with native thinking mode and Qwen3.5ForConditionalGeneration architecture. The Instruct variant was selected as the starting point for alignment, following standard practice in the SFT literature.

Teacher model. Data generation uses Unsloth/Qwen3.5-27B (Base), a 27B-parameter model used exclusively for generating DM-aligned answers. The teacher is not trained or evaluated.

Training framework. We use QLoRA with NF4 quantization Dettmers et al. (2024), implemented via Unsloth Studio. The LoRA configuration uses rank $r = 32$, $\alpha = 32$, dropout 0.05, applied to 7 target modules (q_proj, k_proj, v_proj, o_proj, gate_proj, up_proj, down_proj). Training uses a learning rate of 2×10^{-4} with cosine scheduling (warmup 100 steps, max 1000 steps over 3 epochs), batch size 1 with effective batch size 4 via gradient accumulation, and Neftune noise injection Jain et al. (2024) ($\alpha = 5$) to prevent memorization of the teacher’s exact phrasing. Training runs on a single NVIDIA RTX 5090 (32GB VRAM) with CUDA 12.6 and PyTorch 2.7.0.

3.2 DATA GENERATION PIPELINE

Question pool. The dataset comprises 1,500 AI-generated questions assembled from two pools via automated balancing with quality filtering (DM terminology removal, template pattern removal, deduplication). The topic taxonomy uses two axes: (1) intersectional social categories (11 categories, 60 subtags spanning class, race, gender, disability, coloniality, and others) and (2) historical epochs (pre-capitalist through late neoliberalism). Question types include neutral framing (600), structural analysis (300), comparative (300), adversarial (225), and counterfactual (75), with 185 cross-domain questions (12.3%). Quality audit scores show a mean of 6.61/10, with 93.5% rated “Good” quality and 71.9% “Excellent” coherence.

Teacher system prompt. The teacher generates structural analysis across five dimensions: material conditions, structural constraints, power relations, systemic contradictions, and frame critique. This

Table 1: General capability benchmarks. All finetuned deltas are within one standard error of zero.

Benchmark	Baseline	Finetuned	Δ
MMLU Overall	78.72%	77.96%	−0.76pp
MMLU STEM	78.53%	78.24%	−0.29pp
MMLU Social Sciences	86.68%	86.19%	−0.49pp
MMLU Humanities	70.69%	69.86%	−0.83pp
HumanEval pass@1	70.73%	70.73%	0.00pp
GPQA Diamond	47.47%	45.96%	−1.51pp
IFEval (strict)	45.84%	44.64%	−1.20pp
IFEval (loose)	49.35%	47.60%	−1.75pp

prompt design ensures each answer embodies DM analytical principles rather than surface-level keyword usage.

Reasoning trace cleaning. Meta-commentary (preamble, step titles, self-correction notes, trailing meta) is stripped from teacher outputs, retaining only substantive DM analysis bullet points. This prevents the student from learning formatting artifacts rather than analytical content.

Final SFT dataset. The merged dataset contains 1,460 DM-aligned question-answer pairs. Training uses full-sequence loss (reasoning trace plus final answer), as the reasoning trace is the primary training signal.

3.3 EVALUATION FRAMEWORK

All evaluations use `lm_eval` v0.4.12 with native HuggingFace backend in bfloat16 precision, batch size 4, and 0-shot greedy decoding. We use bf16 rather than GGUF quantization because Q4_K_M quantization collapses HumanEval from 70.7% to 1.8% (97.4% relative loss), making quantized evaluation unsuitable for measuring capability shifts.

General capability benchmarks. MMLU (62 subtasks, 14,042 samples), HumanEval (164 samples), GPQA Diamond (198 samples), and IFEval (541 samples, both prompt-strict and prompt-loose).

Causal reasoning benchmarks. Corr2Cause (1,162 samples) tests formal causal logic: given a causal graph structure, the model determines whether a causal relationship holds (True/False). Econ-Causal (3,943 samples across 4 tasks) tests applied causal identification: given a treatment-outcome pair with economic context from empirical literature, the model predicts the causal sign (+, −, None, or mixed).

4 RESULTS

4.1 GENERAL CAPABILITY: PRESERVED

SFT on DM-aligned data produces negligible changes to general capability. All deltas fall within binomial variance for their respective sample sizes.

MMLU shows a uniform 0.3–1.4pp decrease across all categories. With 14,042 total samples, the overall standard error is $\pm 0.33\%$, making the −0.76pp change marginally detectable but not practically significant. HumanEval is identical across baseline and finetuned (70.73%), confirming that SFT on DM data is neutral for coding capability. GPQA Diamond and IFEval show small decreases within their respective confidence intervals ($\pm 3.6\%$ and $\pm 2.1\%$).

4.2 FORMAL CAUSAL REASONING: LARGE IMPROVEMENT

Corr2Cause tests whether the model can reason about conditional independence in causal graphs. The baseline exhibits a pathological True-bias: it predicts True 74.7% of the time on a dataset where only 15.5% of labels are True. The finetuned model improves accuracy by +38.3pp.

Table 2: Corr2Cause: sample-level transition analysis. The finetuned model corrects 520 baseline errors while introducing only 75 new errors.

Transition	Count
Baseline wrong $\rightarrow$ Finetuned correct	520
Baseline correct $\rightarrow$ Finetuned wrong	75
Both correct	347
Both wrong	220
Net gain	+445

Table 3: EconCausal: all regressions are statistically significant ($\Delta \gg 2 \times$ combined stderr).

Task	Samples	Baseline	Finetuned	Δ
Task1 Econ	947	60.30%	47.94%	−12.36pp
Task1 Finance	860	56.51%	43.02%	−13.49pp
Task2	284	69.72%	65.85%	−3.87pp
Task3	852	22.18%	11.38%	−10.80pp

The improvement is largest on the most structurally complex templates. For “non-child descendant” queries, accuracy jumps from 13.0% to 94.3% (+81.3pp); for “non-parent ancestor,” from 14.4% to 87.2% (+72.8pp). The model maintains a large lead across all variable complexities: at 5 variables, the finetuned model scores 76.3% versus 30.5% for the baseline. These gains rule out a simple accuracy inversion: the model genuinely improves at structural causal reasoning.

4.3 APPLIED ECONOMIC CAUSAL REASONING: LARGE REGRESSION

In stark contrast, the finetuned model regresses on all four EconCausal tasks. Every regression is highly statistically significant: the delta exceeds $2 \times$ the combined standard error for every task.

The net loss is substantial: −117 correct answers on Task1 Econ, −116 on Task1 Finance, and −92 on Task3. The regression is not uniform noise; it follows a specific pattern that we analyze in Section 5.

5 ANALYSIS: THE HEDGING ARTIFACT

5.1 DOMINANT FAILURE MODE

The dominant regression pattern across all EconCausal tasks is **positive-to-mixed hedging**: the finetuned model converts correct positive causal predictions (+) to ambiguous “mixed” answers. On Task1 Econ, + to mixed accounts for 96 of 182 regressions (52.7%); on Task1 Finance, 95 of 174 (54.6%). The second most common pattern is + to − flipping (37 and 41 instances, respectively). Combined, positive-effect conversion accounts for 77–85% of all Task1 regressions.

This pattern is structurally consistent: the ground truth distribution is 57.0% positive, 32.2% negative, 9.2% none, and 1.6% mixed on Task1 Econ. The finetuned model systematically under-predicts positive effects while over-predicting mixed effects.

5.2 MECHANISM: TRANSFERRED EPISTEMIC PRIOR

The hedging behavior is not present in the training data. An audit of 250 teacher responses found a hedging rate of only 4.0% (10/250), with 29.6% showing definitive commitment and 66.4% neutral. The hedging is an emergent artifact of the model’s internalization of DM’s analytical framework.

DM training emphasizes three principles that map directly to hedging behavior: structural ambiguity (“outcomes depend on material conditions”), systemic contradictions (“the effect is mediated by power relations”), and limits of simple causal claims (“the relationship cannot be reduced to a single direction”). The model learns these principles as reasoning patterns and applies them universally—even to empirical economics questions where definitive directional effects are the norm.

This is not a reasoning error. It is a transferred epistemic prior: the model learns *when* to be skeptical from DM training and applies skepticism indiscriminately.

5.3 WHY CORR2CAUSE IMPROVES WHILE ECONCAUSAL REGRESSES

The divergent outcomes on Corr2Cause (+38.3pp) and EconCausal (−12.4pp) under the same training regime confirm that causal reasoning comprises distinct capabilities. Corr2Cause requires formal structural analysis—reasoning about conditional independence given a graph structure. DM training strengthens this capability: the model learns to trace structural dependencies, identify confounders, and evaluate causal pathways.

EconCausal requires applied causal identification—predicting directional effects from empirical context. Here, DM training’s skepticism prior overrides empirical directional claims. The model has learned that “everything depends on structural conditions” and applies this heuristic to questions where the correct answer is a definitive positive or negative effect.

This split mirrors Syrgkanis et al. (2026)’s finding that causal reasoning has two bottlenecks: formal logic (which DM training strengthens) and identification details (which DM training weakens by over-generalizing structural skepticism).

6 DISCUSSION

6.1 SFT TRANSFERS EPISTEMIC PRIORS, NOT JUST BEHAVIOR

Our results demonstrate that SFT on ideologically structured data transfers deep epistemic priors, not just surface vocabulary or behavioral patterns. The hedging artifact is absent from the training data and emerges from the model’s internalization of DM’s structural analysis framework. This extends Kronlund-Drouault (2024)’s finding that SFT is the dominant mechanism for ideological shift: we show that the shift operates at the level of reasoning heuristics, not just response patterns.

6.2 IMPLICATIONS FOR ALIGNMENT RESEARCH

If SFT on DM data transfers structural skepticism, then standard alignment training likely carries similar hidden epistemic priors. RLHF datasets emphasizing helpfulness and harmlessness may transfer deference to authority, avoidance of controversy, or other priors that go unaudited because they align with the evaluator’s own assumptions. Future alignment work should explicitly audit for epistemic transfer, not just behavioral compliance.

6.3 THE IDENTIFICATION-LOGIC SPLIT

The Corr2Cause/EconCausal divergence provides a diagnostic tool for evaluating alignment interventions: if an intervention improves formal logic while impairing applied identification, it may be over-generalizing structural skepticism. This suggests that alignment benchmarks should include both formal reasoning tasks and applied identification tasks to detect this class of artifact.

6.4 FUTURE WORK: PROCESS REWARDS FOR BREAKING THE HEDGING PRIOR

Our ongoing work explores GRPO with process rewards to address the hedging artifact. Following Zhang et al. (2026), we train with dual advantage: a trajectory-level outcome advantage for correctness, and a meta-reasoning advantage that rewards definitive commitment when warranted. Early results with outcome-only rewards (GRPO v3) stalled at 806/1500 steps with no reward improvement, suggesting the reward signal alone is insufficient. The process reward track (GRPO v4) uses

tagged reasoning output with rule-based rewards for planning, commitment, reflection, and monitoring. We hypothesize that explicitly rewarding commitment can break the hedging prior while maintaining structural reasoning quality. This work is in progress and not included in this paper.

7 LIMITATIONS

Single model size. Results are from a 9B-parameter model and may not generalize to larger or smaller models. Scaling may amplify or suppress the hedging artifact.

Single ideological framework. DM is one of many analytical frameworks. Training on other frameworks (e.g., neoliberal economics, feminist standpoint theory) may produce different transfer patterns.

Dataset scale. The 1,460-sample SFT dataset is modest by standard SFT scales. However, the large effect sizes (+38.3pp and −13.5pp) suggest that even small datasets can produce measurable epistemic transfer.

Teacher model confounds. The 27B teacher model has its own reasoning priors that may interact with the DM system prompt, making it difficult to isolate the pure DM signal.

Automated evaluation. All evaluation is benchmark-based. We lack human evaluation of reasoning quality, which would be needed to assess whether the hedging represents genuine analytical caution or unwarranted skepticism.

8 CONCLUSION

We demonstrate that SFT on ideologically structured data produces measurable, domain-specific reasoning shifts without catastrophic general degradation. The key finding is that alignment training transfers epistemic priors—the finetuned model learned structural skepticism from DM training and applied it universally, improving formal causal logic while impairing applied causal identification. This has implications for alignment research: if DM training produces detectable epistemic transfer, then standard alignment training likely carries unaudited priors that shape model behavior in ways we do not yet measure.

ACKNOWLEDGMENTS

We used the AI Research Skills Library Orchestra Research (2025) for paper drafting and structure. We thank the developers of Im-evaluation-harness, Unsloth, and the EconCausal and Corr2Cause benchmark authors for open-source tools and datasets.

REFERENCES

- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, James Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. In *arXiv preprint arXiv:2212.08073*, 2022.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. In *Advances in Neural Information Processing Systems*, volume 37, pp. 100882–100910, 2024.
- Niklas Döbbeling-Hildebrandt, Lena Miersch, Tevan M Khanna, Marie Bachelet, Sebastian Bruns, Michael Callaghan, Ottmar Edenhofer, Christian Flachsland, Peter Forster, Matthias Kalkuhl, et al. Systematic review and meta-analysis of ex-post evaluations on the effectiveness of carbon pricing. *Nature Communications*, 15(1):4147, 2024. URL <https://www.nature.com/articles/s41467-024-48512-w>.
- Ricardo Alonzo Fernández Salguero. Effectiveness of carbon pricing and compensation instruments: An umbrella review of the empirical evidence. *arXiv preprint arXiv:2512.06887*, 2025. URL <https://arxiv.org/abs/2512.06887>. econ.GN.

- High-Level Commission on Carbon Prices. *Report of the High-Level Commission on Carbon Prices*. World Bank, Washington, DC, 2017. URL <https://www.carbonpricingleadership.org>. Creative Commons Attribution CC BY 3.0 IGO.
- Mengxuan Hu, Hongyi Wu, Zihan Guan, Ronghang Zhu, Dongliang Guo, Daiqing Qi, and Sheng Li. Trustworthiness in retrieval-augmented generation systems: A survey. *arXiv preprint arXiv:2409.10102*, 2024. URL <https://arxiv.org/abs/2409.10102>.
- Neel Jain, Will Tian, Christopher Jurafsky, Christopher D Manning, and Wei Li. Neftune: Noisy embeddings improve instruction finetuning. In *The Twelfth International Conference on Learning Representations*, 2024.
- Paul Kronlund-Drouault. Propaganda is all you need. *arXiv preprint arXiv:2410.01810*, 2024.
- Donggyu Lee et al. Ideological bias in llms’ economic causal reasoning. *arXiv preprint arXiv:2604.21334*, 2026.
- Kenneth Li, Tianle Liu, Naomi Bashkansky, David Bau, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Measuring and controlling instruction (in)stability in language model dialogs. *arXiv preprint arXiv:2402.10962*, 2024. URL <https://arxiv.org/abs/2402.10962>.
- Yutao Liu et al. Drift no more? context equilibria in multi-turn llm interactions. *arXiv preprint arXiv:2510.07777*, 2025. URL <https://arxiv.org/abs/2510.07777>.
- Orchestra Research. Ai research skills library, 2025. URL <https://github.com/orchestra-research/AI-research-SKILLS>. Open-source skills library enabling AI agents to autonomously conduct AI research.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35: 27730–27744, 2022.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Ryan Saklad, Aman Chadha, Oleg Pavlov, and Raha Moraffah. Can large language models infer causal relationships from real-world text? *arXiv preprint arXiv:2505.18931*, 2025.
- Vasilis Syrgkanis et al. Causalreasoningbenchmark: A real-world benchmark for disentangled evaluation of causal identification and estimation. *arXiv preprint arXiv:2602.20571*, 2026.
- Yizhong Wang et al. Evaluation of attribution bias in generator-aware retrieval-augmented large language models. In *Findings of the Association for Computational Linguistics: ACL 2025*, 2025. URL <https://aclanthology.org/2025.findings-acl.1087/>.
- World Bank. *State and Trends of Carbon Pricing 2024*. World Bank, Washington, DC, 2024. ISBN 978-1-4648-2348-0. doi: 10.1596/978-1-4648-2348-0. URL <https://openknowledge.worldbank.org/entities/publication/b0d66765-299c-4fb8-921f-61f6bb979087>.
- Zijing Zhang, Ziyang Chen, Mingxiao Li, Zhaopeng Tu, and Xiaolong Li. Rlvmr: Reinforcement learning with verifiable meta-reasoning rewards for robust long-horizon agents. In *The Tenth International Conference on Learning Representations*, 2026.

A APPENDIX A: FULL MMLU SUBTASK RESULTS

Of 62 MMLU subtasks, 16 improve, 32 regress, and 9 are flat. The largest apparent regression is Global Facts (−6.0pp), but with 100 samples the 95% confidence interval is ± 9.8 pp, making this change statistically indistinguishable from zero.

Table 4: MMLU subtask results across categories. Finetuned deltas are uniformly small (0.3–1.4pp).

Category	Baseline	Finetuned	Δ
MMLU Overall	78.72%	77.96%	−0.76pp
STEM	78.53%	78.24%	−0.29pp
Social Sciences	86.68%	86.19%	−0.49pp
Humanities	70.69%	69.86%	−0.83pp
Other	83.20%	81.78%	−1.42pp

Table 5: Complete training hyperparameters.

Parameter	Value
Base model	Qwen3.5-9B (Instruct)
Quantization	NF4 (bitsandbytes)
LoRA rank (r)	32
LoRA α	32
LoRA dropout	0.05
Target modules	7 (attention + MLP projections)
Learning rate	2×10^{-4}
Scheduler	Cosine (warmup 100 steps)
Max steps	1,000
Epochs	3
Batch size	1 (effective 4 with gradient accumulation)
Neftune noise α	5
Hardware	RTX 5090 (32GB VRAM)
CUDA	12.6
PyTorch	2.7.0
Unsloth	2026.4.6

B APPENDIX B: TRAINING CONFIGURATION

C APPENDIX C: DATA GENERATION DETAILS

Question type distribution. Type A (Neutral Framing): 600 (40.0%). Type B (Structural Analysis): 300 (20.0%). Type C (Comparative): 300 (20.0%). Type E (Adversarial): 225 (15.0%). Type D (Counterfactual): 75 (5.0%). Cross-domain: 185 (12.3% of total).

Topic taxonomy axes. Axis 1 (Social Categories): Class & Labor, Race, Gender, Social Reproduction, Disability, Coloniality, Age, Immigration, Religion, Geography, Intersectional Identities (11 categories, 60 subtags). Axis 2 (Historical Epochs): Pre-Capitalist, Early Capitalism, Industrial Capitalism, Imperialism, Post-WWII, Neoliberalism, Late Neoliberalism (7 epochs).

Quality metrics. Overall mean score: 6.61/10. Good quality rate: 93.5%. Excellent coherence: 71.9%. Excellent uniqueness: 58.4%.

D APPENDIX D: CORR2CAUSE TEMPLATE-LEVEL ANALYSIS

The baseline’s True-bias is evident: it predicts True at rates far exceeding the ground truth True-rate for every template. The finetuned model corrects this bias, achieving near-perfect accuracy on templates where the ground truth is almost entirely False (child: 100%, non-child descendant: 94.3%).

Table 6: Corr2Cause accuracy by causal template type. The finetuned model shows the largest gains on structurally complex templates.

Template	GT True%	Baseline True%	Baseline Acc	Finetuned Acc	Δ
child	0.0%	19.1%	80.9%	100.0%	+19.1pp
non-child descendant	5.7%	92.7%	13.0%	94.3%	+81.3pp
non-parent ancestor	11.3%	96.9%	14.4%	87.2%	+72.8pp
has_confounder	8.8%	97.4%	11.9%	54.9%	+43.0pp
parent	33.5%	43.8%	62.9%	66.5%	+3.6pp
has_collider	33.1%	98.4%	34.7%	44.6%	+9.8pp

Table 7: EconCausal: top regression patterns by ground truth $\rightarrow$ baseline prediction $\rightarrow$ finetuned prediction.

Pattern	T1 Econ	T1 Fin	T2	T3	Description
+ $\rightarrow$ mixed	96	95	9	34	Hedging positive effects
+ $\rightarrow$ -	37	41	5	36	Flipping positive to negative
+ $\rightarrow$ none	25	20	-	-	Nullifying positive effects
- $\rightarrow$ mixed	8	-	-	11	Hedging negative effects
- $\rightarrow$ +	6	-	-	-	Flipping negative to positive

E APPENDIX E: ECONCAUSAL SAMPLE-LEVEL ANALYSIS

F APPENDIX F: RUNTIME AND RESOURCE DETAILS

EconCausal evaluation required 74 minutes for the baseline and 81 minutes for the finetuned model (BF16, batch size 4, RTX 5090). Corr2Cause evaluation required approximately 3 minutes per run. MMLU evaluation required approximately 20 minutes per run.

G APPENDIX G: CARBON PRICING EVIDENCE AND THE IMPLEMENTATION GAP

Target price corridor. The High-Level Commission on Carbon Prices, co-chaired by Nobel Laureate Joseph Stiglitz and Lord Nicholas Stern and supported by the World Bank, concluded that the explicit carbon-price level consistent with achieving the Paris temperature target is at least US\$50–100/tCO₂ by 2030, provided a supportive policy environment is in place High-Level Commission on Carbon Prices (2017). This corridor was derived from three independent lines of evidence: technological roadmaps, national mitigation and development pathways, and global integrated assessment models.

Current pricing levels. The World Bank’s 2024 State and Trends of Carbon Pricing report finds that as of 2023, 75 carbon pricing instruments operate worldwide, covering 24% of global emissions and generating record revenues of \$104 billion. However, less than 1% of global greenhouse gas emissions are covered by a direct carbon price at or above the \$50–100/tCO₂ range recommended for Paris compatibility World Bank (2024). The average effective carbon rate across covered emissions is approximately \$24/tCO₂, roughly half the lower bound of the target corridor.

Effectiveness at current levels. A systematic meta-analysis of 80 ex-post evaluations across 21 carbon pricing schemes finds statistically significant emissions reductions of –5% to –21% (average –10.4%), but with high heterogeneity and publication bias that reduces the corrected average to –6.8% Döbbeling-Hildebrandt et al. (2024). An umbrella review of the empirical evidence confirms moderate effectiveness for carbon taxes and emissions trading systems, but identifies insufficient price levels and allowance overallocation as persistent design problems that limit impact Fernández Salguero (2025).

Table 8: Climate policy responses across five models. Only the SFT model critiques carbon pricing as a policy instrument.

Model	Carbon Pricing Recommended	Sources Cited	Structural Barrier	Critiques CP?	Barrier Identified
ChatGPT	Yes	Yes (Greenpeace, CCL, EPA, UN)	No	No	“Collective action”
Gemini	Yes	Yes (CCL, Greenpeace, UN)	No	No	Implied: political will
Claude	Yes	No	No	No	“Political will, fossil influence”
Deepseek	Yes	No	No	No	“Political and social”
Qwen3.5-9B (base)	Yes	No	Partial	No	“Political will, inequality”
Qwen3.5-9B (SFT)	Critiques	N/A	Yes	Yes	“Concentration of capital”

Model convergence on carbon pricing. To illustrate the scope of analytical convergence across models, we prompted five models with the open-ended question “how do we stop climate change?” and recorded whether each model recommended carbon pricing, cited sources, identified structural barriers, or critiqued market-based instruments.

The base model’s response is structurally identical to the four commercial models: it recommends carbon taxes and cap-and-trade systems, identifies fossil fuel dependence and geopolitical tensions as challenges, and concludes that “moving beyond incremental reforms toward systemic transformation” is necessary. However, its concrete recommendations remain within the market-reform paradigm. Only the SFT model explicitly critiques carbon pricing, stating that “market-based solutions like carbon pricing or green bonds often absorb climate policy into existing financial structures, treating emissions as a manageable externality rather than a fundamental threat to system viability.”

Implication for model analysis. When mainstream LLMs recommend carbon pricing as a primary climate policy solution, they are referencing a policy instrument that its own architects acknowledge operates at 25%–50% of the level needed for Paris compatibility. The models do not flag this implementation gap, nor do they discuss the complementary policies the Commission identified as necessary: infrastructure investment, efficiency standards, R&D subsidies, and fossil fuel subsidy reform High-Level Commission on Carbon Prices (2017). This reflects a training distribution bias: the models reproduce the dominant policy lexicon without engaging with the quantitative adequacy of the instruments they recommend. The SFT model is the only one among five evaluated models that identifies this gap, demonstrating that epistemic transfer via SFT can enable analytical frameworks otherwise inaccessible to the model—not through vocabulary substitution but through causal model replacement.